Neural Networks

Part 7a: representation and forward propagation. Part 7b: backpropagation, deep networks and training. Part 7c: the mathematics behind backpropagation

Khiem Nguyen

Email	khiem.nguyen@glasgow.ac.uk
MS Teams	khiem.nguyen@glasgow.ac.uk
WhatsApp	+44 7729 532071 (emergency only)

March 2026



This lecture

A neural network is built from many logistic-regression-like units. This lecture shows how it computes its output, how it is trained, and why the training algorithm works.

Part 7a: representation

- neurons, layers, weights $\mathbf{W}^{[l]}$ and biases $\mathbf{b}^{[l]}$;
- **forward propagation** for a batch, as in PyTorch;
- **activation functions**, and why they must be nonlinear.

Part 7b: training

- computational graphs and the **chain rule**;
- **backpropagation** for one hidden layer and for L layers;
- initialisation, **vanishing gradients**, training in **PyTorch**.

Part 7c: the mathematics

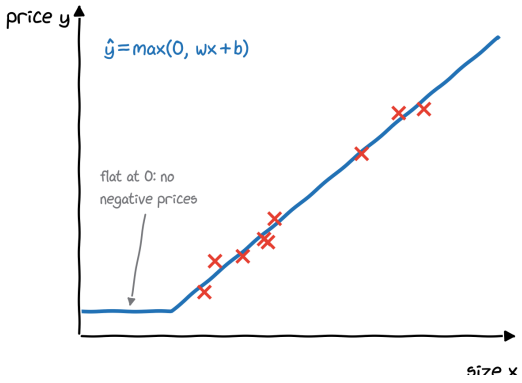
- the chain rule for vectors and matrices;
- "easy" local and accumulated derivatives;
- a **worked example** by hand, checked numerically.

An optional appendix proves backpropagation with matrix calculus.

Part 7a: representation and forward propagation

From linear models to neural networks

The simplest neural network: one neuron



Predict the price y of a house from its size x. A straight line would predict negative prices for small houses, so cut it off at zero:

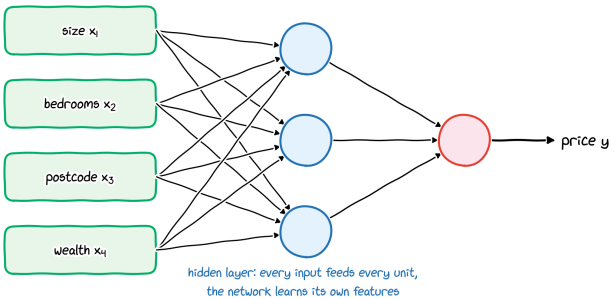
$$\hat{y} = \mathrm{ReLU}(wx + b),$$

$$\mathrm{ReLU}(z) = \max(0, z) = \begin{cases} z & z \geq 0 \\ 0 & z < 0 \end{cases}$$

This is a single **neuron** (or **unit**):

1. compute a weighted sum $z = wx + b$, as in linear regression;
2. apply an **activation function**, here ReLU, to get the output $a = \mathrm{ReLU}(z)$.

From one neuron to a hidden layer



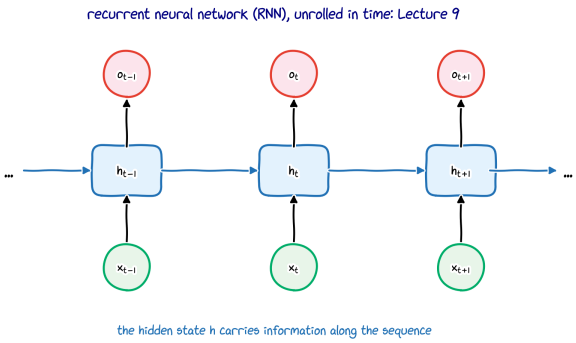
In practice we do not design these features by hand: every input feeds every unit of the **hidden layer**, and training decides what each unit computes.

Supervised learning with neural networks

input $\mathbf{x}$	output y	application	typical network
house features	price	real estate	fully connected
ad and user information	click on ad? (0/1)	online advertising	fully connected
image	object class $(1, \dots, 1000)$	photo tagging	convolutional (CNN)
audio	text transcript	speech recognition	recurrent (RNN), transformer
English text	Chinese text	machine translation	transformer, RNN
image, radar	positions of other cars	autonomous driving	custom / hybrid

- **Structured data** (tables of features) and **unstructured data** (images, audio, text) can both be handled by neural networks.
- Unstructured data is where neural networks have had their biggest successes.

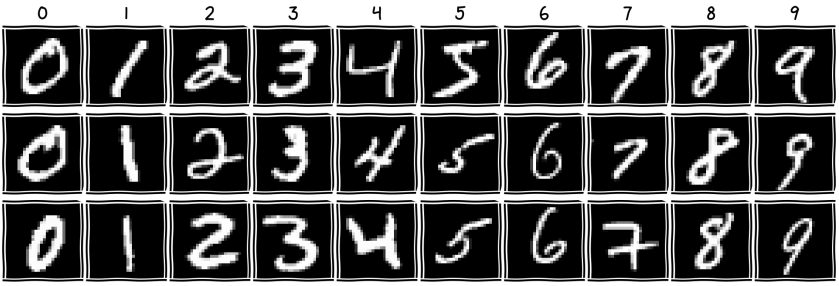
Network architectures in this course



All three are trained in the same way: a loss function, gradients by backpropagation, and gradient descent.

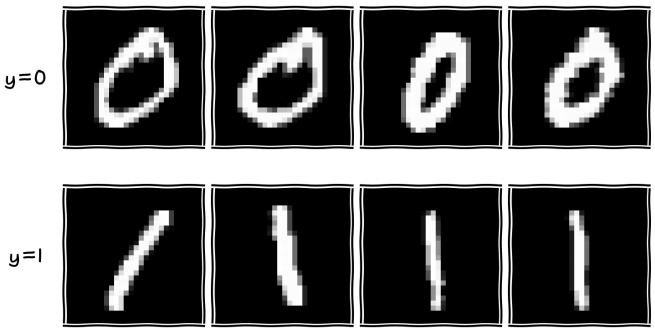
A first problem: handwritten digits

The MNIST digits



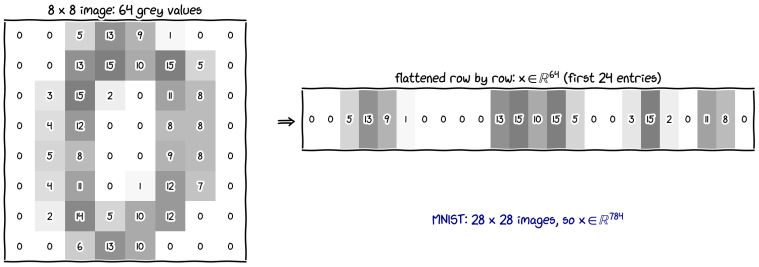
- 70,000 images of handwritten digits 0 to 9: 60,000 for training and 10,000 for testing.
- Each image is greyscale with 28×28 pixels; each pixel is an integer from 0 (background) to 255 (full ink).
- A classic benchmark for image classification.

Start simple: 0 versus 1



A **binary** classification problem, as in Lecture 3a: is the digit a 0 ($y = 0$) or a 1 ($y = 1$)? We will later extend it to all ten digits with softmax (Lecture 3b).

Images as feature vectors



- Stack the rows of pixels into one long vector: an 8×8 image becomes $\mathbf{x} \in \mathbb{R}^{64}$, an MNIST image $\mathbf{x} \in \mathbb{R}^{784}$.
- In NumPy: `x = image.flatten()`; in PyTorch: `x = image.reshape(-1)` or the layer `nn.Flatten()`.
- The network sees 784 numbers. Fully connected networks ignore that neighbouring pixels belong together; convolutional networks (Lecture 8) use this structure.

Notation for the data

- One training example: $(\mathbf{x}, y)$ with $\mathbf{x} \in \mathbb{R}^d$ (d features, e.g. $d = 784$) and $y \in \{0, 1\}$.
- Training set: m examples $\{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(m)}, y^{(m)})\}$.

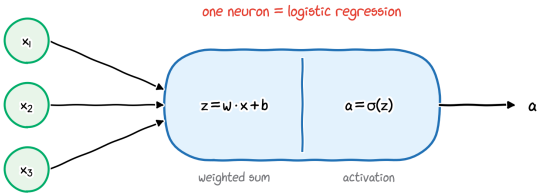
As in Lecture 2b, the examples are the **rows** of the data matrix:

$$\mathbf{X} = \begin{pmatrix} — \ \mathbf{x}^{(1)} \ — \\ — \ \mathbf{x}^{(2)} \ — \\ \vdots \\ — \ \mathbf{x}^{(m)} \ — \end{pmatrix} \in \mathbb{R}^{m \times d}, \qquad \mathbf{y} = \begin{pmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{pmatrix} \in \mathbb{R}^m.$$

Many textbooks store the examples as **columns** ($\mathbf{X} \in \mathbb{R}^{d \times m}$). We use rows because PyTorch, NumPy and scikit-learn do; the formulas then contain transposes in different places, but they describe the same network.

Neural network representation

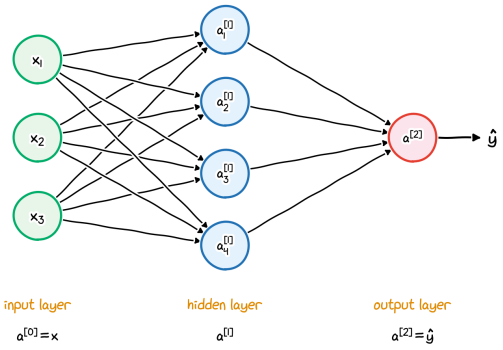
Logistic regression is a one-neuron network



$$z = \mathbf{w} \cdot \mathbf{x} + b, \qquad a = \sigma(z) = \frac{1}{1 + e^{-z}}, \qquad \hat{y} = a$$

- From now on we write σ for the sigmoid (it was g in Lecture 3a), and g for a general activation function.
- Trained by minimising the binary cross-entropy (Lecture 3a):
  $J(\mathbf{w}, b) = -\frac{1}{m} \sum_{i=1}^{m} \left[y^{(i)} \log a^{(i)} + (1 - y^{(i)}) \log(1 - a^{(i)}) \right]$.

A network with one hidden layer

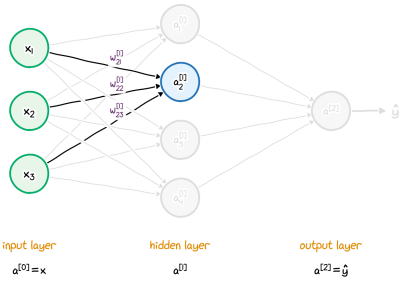


- Each hidden unit $a_j^{[1]}$ is a function of (x_1, x_2, x_3).
- The output $a^{[2]}$ is a function of $(a_1^{[1]}, \ldots, a_4^{[1]})$...
- ...so $\hat{y} = a^{[2]}$ is a function of $\mathbf{x}$: the network is just a (complicated) model $f(\mathbf{x})$.

Notation

- Superscript $[l]$: layer number. $\mathbf{a}^{[0]} = \mathbf{x}$ is the input.
- Subscript j: unit number within the layer.
- $n^{[l]}$: number of units in layer l; here $n^{[0]} = 3$, $n^{[1]} = 4$, $n^{[2]} = 1$.
- This is a **two-layer** network: the input layer is not counted.

Computing the hidden layer, one unit at a time



Each hidden unit is a logistic-regression unit with its **own** weights and bias:

$$z_1^{[1]} = \mathbf{w}_1^{[1]} \cdot \mathbf{a}^{[0]} + b_1^{[1]}, \quad a_1^{[1]} = \sigma(z_1^{[1]})$$
$$z_2^{[1]} = \mathbf{w}_2^{[1]} \cdot \mathbf{a}^{[0]} + b_2^{[1]}, \quad a_2^{[1]} = \sigma(z_2^{[1]})$$
$$z_3^{[1]} = \mathbf{w}_3^{[1]} \cdot \mathbf{a}^{[0]} + b_3^{[1]}, \quad a_3^{[1]} = \sigma(z_3^{[1]})$$
$$z_4^{[1]} = \mathbf{w}_4^{[1]} \cdot \mathbf{a}^{[0]} + b_4^{[1]}, \quad a_4^{[1]} = \sigma(z_4^{[1]})$$

with $\mathbf{w}_j^{[1]} = (w_{j1}^{[1]}, w_{j2}^{[1]}, w_{j3}^{[1]})$: $w_{jk}^{[1]}$ is the weight **from** input k **to** unit j.

Matrix form

Collect the weight vectors as the **rows** of a matrix, and the biases in a vector:

$$\mathbf{W}^{[1]} = \begin{pmatrix} w_{11}^{[1]} & w_{12}^{[1]} & w_{13}^{[1]} \\ w_{21}^{[1]} & w_{22}^{[1]} & w_{23}^{[1]} \\ w_{31}^{[1]} & w_{32}^{[1]} & w_{33}^{[1]} \\ w_{41}^{[1]} & w_{42}^{[1]} & w_{43}^{[1]} \end{pmatrix} = \begin{pmatrix} — & \mathbf{w}_1^{[1]} & — \\ — & \mathbf{w}_2^{[1]} & — \\ — & \mathbf{w}_3^{[1]} & — \\ — & \mathbf{w}_4^{[1]} & — \end{pmatrix} \in \mathbb{R}^{4\times 3}, \qquad \mathbf{b}^{[1]} = \begin{pmatrix} b_1^{[1]} \\ b_2^{[1]} \\ b_3^{[1]} \\ b_4^{[1]} \end{pmatrix} \in \mathbb{R}^4.$$

Then the four equations become one, and the output layer looks the same:

$$\begin{cases} \mathbf{z}^{[1]} = \mathbf{W}^{[1]}\mathbf{a}^{[0]} + \mathbf{b}^{[1]}, & \mathbf{a}^{[1]} = \sigma(\mathbf{z}^{[1]}), & \mathbf{W}^{[1]} \in \mathbb{R}^{4\times 3} \\ z^{[2]} = \mathbf{W}^{[2]}\mathbf{a}^{[1]} + b^{[2]}, & a^{[2]} = \sigma(z^{[2]}), & \mathbf{W}^{[2]} \in \mathbb{R}^{1\times 4} \end{cases}$$

σ is applied to each entry of the vector separately.

Any layer, any size

For layer $l = 1, 2, \ldots, L$ with $n^{[l]}$ units and activation function $g^{[l]}$:

$$\mathbf{z}^{[l]} = \mathbf{W}^{[l]}\mathbf{a}^{[l-1]} + \mathbf{b}^{[l]}, \qquad \mathbf{a}^{[l]} = g^{[l]}(\mathbf{z}^{[l]}), \qquad \mathbf{W}^{[l]} \in \mathbb{R}^{n^{[l]} \times n^{[l-1]}}, \quad \mathbf{b}^{[l]} \in \mathbb{R}^{n^{[l]}}.$$

$$\underbrace{\begin{pmatrix} a_1^{[l]} \\ \vdots \\ a_{n^{[l]}}^{[l]} \end{pmatrix}}_{\mathbf{a}^{[l]}} = g^{[l]} \left[\underbrace{\begin{pmatrix} w_{11}^{[l]} & \cdots & w_{1n^{[l-1]}}^{[l]} \\ \vdots & \ddots & \vdots \\ w_{n^{[l]}1}^{[l]} & \cdots & w_{n^{[l]}n^{[l-1]}}^{[l]} \end{pmatrix}}_{\mathbf{W}^{[l]}\text{: one row per unit}} \underbrace{\begin{pmatrix} a_1^{[l-1]} \\ \vdots \\ a_{n^{[l-1]}}^{[l-1]} \end{pmatrix}}_{\mathbf{a}^{[l-1]}} + \underbrace{\begin{pmatrix} b_1^{[l]} \\ \vdots \\ b_{n^{[l]}}^{[l]} \end{pmatrix}}_{\mathbf{b}^{[l]}} \right]$$

- The output of one layer is the input of the next: $\mathbf{x} = \mathbf{a}^{[0]} \to \mathbf{a}^{[1]} \to \ldots \to \mathbf{a}^{[L]} = \hat{\mathbf{y}}$.
- Computing these in order is called **forward propagation**.

Many examples at once: the form used by PyTorch

Write one example as a **row** vector and stack the m examples as rows, $\mathbf{A}^{[0]} = \mathbf{X} \in \mathbb{R}^{m \times d}$. Transposing $\mathbf{z} = \mathbf{W}\mathbf{a} + \mathbf{b}$ row by row gives:

Forward propagation for a batch (two-layer network)

$$\mathbf{Z}^{[1]} = \mathbf{A}^{[0]}\big(\mathbf{W}^{[1]}\big)^{\top} + \mathbf{b}^{[1]}, \quad \mathbf{A}^{[1]} = \sigma(\mathbf{Z}^{[1]}), \quad \mathbf{Z}^{[1]}, \mathbf{A}^{[1]} \in \mathbb{R}^{m \times n^{[1]}}$$

$$\mathbf{Z}^{[2]} = \mathbf{A}^{[1]}\big(\mathbf{W}^{[2]}\big)^{\top} + \mathbf{b}^{[2]}, \quad \mathbf{A}^{[2]} = \sigma(\mathbf{Z}^{[2]}), \quad \mathbf{Z}^{[2]}, \mathbf{A}^{[2]} \in \mathbb{R}^{m \times 1}$$

- Row i of $\mathbf{A}^{[l]}$ holds the activations of layer l for example i.
- The bias row vector is added to **every** row (broadcasting, Lecture 6).
- Each layer feeds the next: $\mathbf{Z}^{[2]}$ uses $\mathbf{A}^{[1]}$, not $\mathbf{A}^{[0]}$.

Forward propagation in PyTorch

```
import torch
from torch import nn

torch.manual_seed(0)
model = nn.Sequential(
    nn.Linear(3, 4),    # W1: (4, 3), b1: (4,)
    nn.Sigmoid(),
    nn.Linear(4, 1),    # W2: (1, 4), b2: (1,)
    nn.Sigmoid(),
)
W1, b1 = model[0].weight, model[0].bias
W2, b2 = model[2].weight, model[2].bias

X = torch.rand(5, 3)                # m = 5, d = 3
A1 = torch.sigmoid(X @ W1.T + b1) # (5, 4)
A2 = torch.sigmoid(A1 @ W2.T + b2)# (5, 1)
torch.allclose(A2, model(X))      # True
```

- `nn.Linear(in, out)` stores $\mathbf{W}^{[l]}$ with shape (out, in), exactly as on the previous slides, and computes `X @ W.T + b`.
- So `model(X)` **is** the batch formula: our hand-written forward pass gives the same result.
- `nn.Sequential` applies the layers in order: forward propagation.
- The weights start with small random values; training (part 7b) adjusts them.

Counting parameters

Layer l has $n^{[l]} \times n^{[l-1]}$ weights and $n^{[l]}$ biases, i.e. $n^{[l]}(n^{[l-1]} + 1)$ parameters.

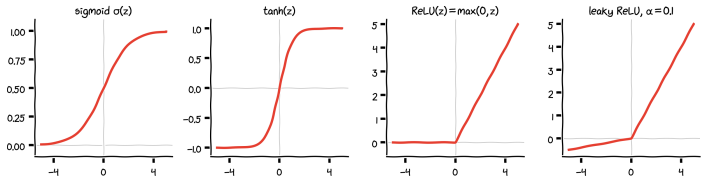
network	layer 1	layer 2	total
$3 \to 4 \to 1$ (previous slide)	$4 \times (3 + 1) = 16$	$1 \times (4 + 1) = 5$	21
MNIST 0 vs 1: $784 \to 25 \to 1$	$25 \times 785 = 19\,625$	$1 \times 26 = 26$	$19\,651$

```
sum(p.numel() for p in model.parameters())      # 21
```

Almost all parameters sit in the first layer, because the input has 784 features. Logistic regression on the same images would have 785.

Activation functions

Common activation functions



$$\sigma(z) = \frac{1}{1+e^{-z}}, \qquad \tanh(z) = \frac{e^{z}-e^{-z}}{e^{z}+e^{-z}},$$

$$\mathrm{ReLU}(z) = \max(0, z), \qquad \mathrm{LeakyReLU}(z) = \begin{cases} z & z \geq 0 \\ \alpha z & z < 0 \end{cases}$$

- Sigmoid: outputs in $(0, 1)$. tanh: outputs in $(-1, 1)$, centred at 0; $\tanh(z) = 2\sigma(2z) - 1$.
- Leaky ReLU: a small slope $\alpha > 0$ for negative inputs (typically 0.01; the figure uses 0.1), so the output is $\alpha z < 0$ there.

Which activation where?

Output layer: chosen by the type of target.

task	output activation	loss (PyTorch)
binary classification	sigmoid: $P(y = 1 \mid \mathbf{x})$	`nn.BCEWithLogitsLoss`
K classes	softmax (Lecture 3b)	`nn.CrossEntropyLoss`
regression	none (identity)	`nn.MSELoss`

Hidden layers: usually **ReLU**, the default choice in modern networks.

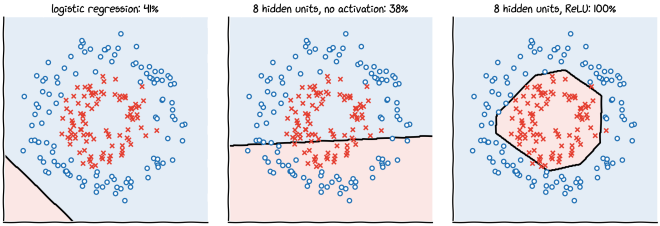
- It is cheap to compute, and its slope is 1 for all positive inputs.
- Sigmoid and tanh **saturate**: for large $|z|$ their slope is almost 0, which slows down gradient-based training (part 7b).
- A ReLU unit whose input is always negative outputs 0 and stops learning ("dead" unit); leaky ReLU avoids this.

Why must activations be nonlinear?

Suppose every activation is the identity, $g(z) = z$. Then

$$\mathbf{a}^{[2]} = \mathbf{W}^{[2]}\left(\mathbf{W}^{[1]}\mathbf{x} + \mathbf{b}^{[1]}\right) + \mathbf{b}^{[2]} = \underbrace{\mathbf{W}^{[2]}\mathbf{W}^{[1]}}_{\mathbf{W}'}\,\mathbf{x} + \underbrace{\mathbf{W}^{[2]}\mathbf{b}^{[1]} + \mathbf{b}^{[2]}}_{\mathbf{b}'}:$$

a **linear** model again, however many layers we stack.



Points inside a ring cannot be separated by a straight line: logistic regression and a network without activation both fail (test accuracy 41% and 38%). With ReLU, the same 8 hidden units give 100%.

Summary of part 7a

- A **neuron** computes $a = g(\mathbf{w} \cdot \mathbf{x} + b)$; logistic regression is a one-neuron network.
- A **layer** of $n^{[l]}$ neurons: $\mathbf{z}^{[l]} = \mathbf{W}^{[l]}\mathbf{a}^{[l-1]} + \mathbf{b}^{[l]}$, $\mathbf{a}^{[l]} = g^{[l]}(\mathbf{z}^{[l]})$, with $\mathbf{W}^{[l]} \in \mathbb{R}^{n^{[l]} \times n^{[l-1]}}$.
- **Forward propagation** for a batch (rows = examples, as in PyTorch): $\mathbf{Z}^{[l]} = \mathbf{A}^{[l-1]}(\mathbf{W}^{[l]})^\top + \mathbf{b}^{[l]}$, $\mathbf{A}^{[l]} = g^{[l]}(\mathbf{Z}^{[l]})$, starting from $\mathbf{A}^{[0]} = \mathbf{X}$.
- **Parameters**: $n^{[l]}(n^{[l-1]} + 1)$ per layer.
- **Activations**: ReLU in hidden layers; sigmoid, softmax or none at the output. They must be **nonlinear**, otherwise the whole network is linear.
- In PyTorch: `nn.Sequential(nn.Linear(d, n1), nn.ReLU(), nn.Linear(n1, 1), ...)`.

Next, part 7b: how to train a network: the computational graph, backpropagation and gradient descent, deeper networks, and MNIST in PyTorch.

Check your understanding: part 7a

1. A network has layers of sizes $10 \to 8 \to 8 \to 3$. What are the shapes of $\mathbf{W}^{[1]}$, $\mathbf{W}^{[2]}$ and $\mathbf{W}^{[3]}$? How many parameters does it have?
2. For a batch of $m = 32$ examples, what is the shape of $\mathbf{Z}^{[1]}$ in that network?
3. In `nn.Linear(784, 25)`, what is the shape of `.weight`?
4. Why is a network with 10 hidden layers but no nonlinear activation no more powerful than logistic regression?
5. Which output activation would you use to predict (a) a house price, (b) one of ten digits, (c) spam or not spam?

Part 7b: backpropagation, deep networks and training

Where we are

In part 7a we computed the output of a network, **forward propagation**:

$$\mathbf{Z}^{[l]} = \mathbf{A}^{[l-1]}\big(\mathbf{W}^{[l]}\big)^\top + \mathbf{b}^{[l]}, \qquad \mathbf{A}^{[l]} = g^{[l]}\big(\mathbf{Z}^{[l]}\big), \qquad l = 1, \dots, L, \quad \mathbf{A}^{[0]} = \mathbf{X}.$$

Part 7b **trains** the network, exactly as in Lectures 2 and 3:

- a cost J (cross-entropy for classification), minimised by gradient descent;
- the new difficulty: the gradient with respect to **every** $\mathbf{W}^{[l]}$ and $\mathbf{b}^{[l]}$, through many layers;
- the answer: **backpropagation**, the chain rule applied backwards through the network.

Computational graphs and the chain rule

The chain rule

If J depends on θ only through an intermediate quantity v, i.e. $J = J(v)$ and $v = v(\theta)$, then

$$\frac{\partial J}{\partial \theta} = \frac{\partial J}{\partial v}\,\frac{\partial v}{\partial \theta}.$$

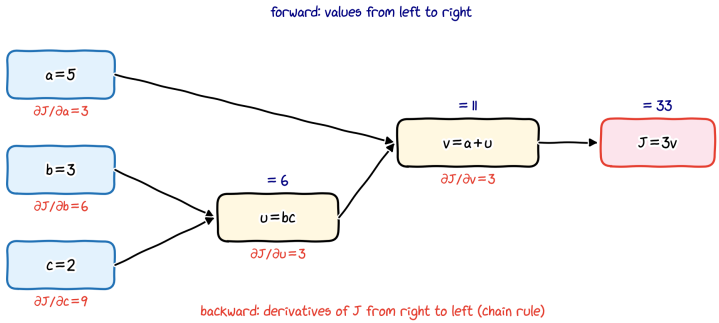
If θ influences J through **several** intermediate quantities $v_1, \dots, v_K$, the contributions add up:

$$\frac{\partial J}{\partial \theta} = \sum_{k=1}^{K} \frac{\partial J}{\partial v_k}\,\frac{\partial v_k}{\partial \theta}.$$

Example. $f = a_1 a_2$ with $a_1 = \theta^2$ and $a_2 = 3\theta$: $\frac{\partial f}{\partial \theta} = a_2 \cdot 2\theta + a_1 \cdot 3 = 9\theta^2$, which agrees with differentiating $f = 3\theta^3$ directly.

Backpropagation is nothing more than applying these two rules systematically, from the output back to the parameters.

A small computational graph



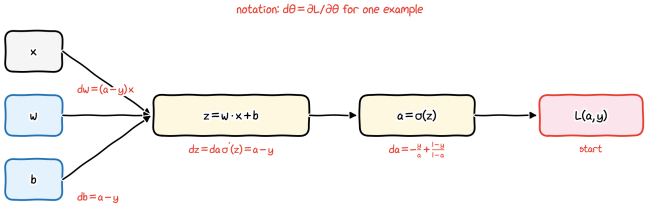
$J(a,b,c) = 3(a+bc)$, evaluated at $a=5$, $b=3$, $c=2$. **Backward**, using values from the forward pass: $\frac{\partial J}{\partial v} = 3$, $\frac{\partial J}{\partial a} = \frac{\partial J}{\partial v} \cdot 1 = 3$, $\frac{\partial J}{\partial u} = 3$, $\frac{\partial J}{\partial b} = \frac{\partial J}{\partial u}\, c = 6$, $\frac{\partial J}{\partial c} = \frac{\partial J}{\partial u}\, b = 9$.

Why go backwards?

- Each derivative $\frac{\partial J}{\partial (\cdot)}$ is computed **once** and then **reused** by every node further back: $\frac{\partial J}{\partial u}$ serves both b and c.
- One backward pass gives the derivatives with respect to **all** parameters, at roughly the cost of one forward pass.
- Each step only needs a **local** derivative ("easy"), such as $\partial u / \partial b = c$, and the derivative passed back from the next node.
- A network with a million parameters still needs only one forward and one backward pass per gradient. This is what `loss.backward()` does in PyTorch (Lecture 6).

Backpropagation for logistic regression

Logistic regression as a graph



For one example, with $z = \mathbf{w} \cdot \mathbf{x} + b$, $a = \sigma(z)$ and $L(a, y) = -y \log a - (1 - y) \log(1 - a)$. We write $d\theta = \partial L / \partial \theta$ (the name of the variable in code):

$$da = -\frac{y}{a} + \frac{1 - y}{1 - a}, \qquad dz = da\ \sigma'(z) = da\ a(1 - a) = a - y,$$
$$d\mathbf{w} = dz\ \mathbf{x} = (a - y)\,\mathbf{x}, \quad db = dz = a - y.$$

Only "easy" local derivatives are needed: $\partial a / \partial z = a(1 - a)$, $\partial z / \partial \mathbf{w} = \mathbf{x}$, $\partial z / \partial b = 1$. Averaging over the examples gives the gradient of Lecture 3a.

Backpropagation for a two-layer network

The forward pass

One example, as a **row** vector (part 7a), hidden activation g, sigmoid output:

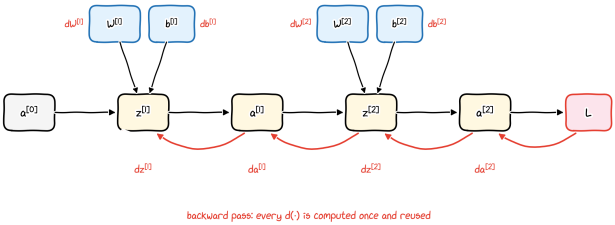
$$\mathbf{z}^{[1]} = \mathbf{a}^{[0]}\big(\mathbf{W}^{[1]}\big)^\top + \mathbf{b}^{[1]}, \quad \mathbf{a}^{[1]} = g(\mathbf{z}^{[1]}), \quad \text{with } \mathbf{a}^{[0]} = \mathbf{x},$$
$$z^{[2]} = \mathbf{a}^{[1]}\big(\mathbf{W}^{[2]}\big)^\top + b^{[2]}, \quad a^{[2]} = \sigma(z^{[2]}), \quad L(a^{[2]}, y) = -y \log a^{[2]} - (1-y)\log\left(1 - a^{[2]}\right).$$

Goal: the gradient of L with respect to all parameters,

$$\frac{\partial L}{\partial \mathbf{W}^{[1]}}, \qquad \frac{\partial L}{\partial \mathbf{b}^{[1]}}, \qquad \frac{\partial L}{\partial \mathbf{W}^{[2]}}, \qquad \frac{\partial L}{\partial b^{[2]}}.$$

The parameters of layer 1 influence L only through $\mathbf{z}^{[1]}$, then $\mathbf{a}^{[1]}$, then $z^{[2]}$ and $a^{[2]}$. So we work backwards along this chain.

The backward pass



$$da^{[2]} \to dz^{[2]} \to \begin{cases} d\mathbf{W}^{[2]},\ db^{[2]} \\ d\mathbf{a}^{[1]} \to d\mathbf{z}^{[1]} \to d\mathbf{W}^{[1]},\ d\mathbf{b}^{[1]} \end{cases}$$

- Layer 2 is handled exactly like logistic regression.
- For layer 1 we continue the **same** process one layer further, starting from $d\mathbf{a}^{[1]}$.
- Each arrow needs only one local derivative: $\partial a/\partial z$, $\partial z/\partial a$ (one layer back), or $\partial z/\partial \theta$ (the parameters).

The backpropagation equations

Two-layer network, one example (rows; $\odot$: element-wise product)

$$
\begin{aligned}
dz^{[2]} &= a^{[2]} - y && && \text{(sigmoid output, cross-entropy loss)} \\
d\mathbf{W}^{[2]} &= \left(dz^{[2]}\right)^{\top} \mathbf{a}^{[1]}, && db^{[2]} = dz^{[2]} && \\
d\mathbf{a}^{[1]} &= dz^{[2]}\, \mathbf{W}^{[2]} && && \text{(pass the error one layer back)} \\
d\mathbf{z}^{[1]} &= d\mathbf{a}^{[1]} \odot g'(\mathbf{z}^{[1]}) && && \text{(through the activation)} \\
d\mathbf{W}^{[1]} &= \left(d\mathbf{z}^{[1]}\right)^{\top} \mathbf{a}^{[0]}, && d\mathbf{b}^{[1]} = d\mathbf{z}^{[1]} &&
\end{aligned}
$$

- **Shapes**: $d\mathbf{W}^{[l]}$ has the same shape as $\mathbf{W}^{[l]}$: $(n^{[l]} \times 1)(1 \times n^{[l-1]})$.
- **Activation derivatives**: sigmoid $g'(z) = a(1-a)$; tanh: $g'(z) = 1 - a^2$; ReLU: $g'(z) = 1$ if $z > 0$, else 0.

Where do these formulas come from? (optional)

Entry j of layer l is $z_j = \sum_k w_{jk}\, a_k + b_j$, where a_k are the activations of layer $l-1$. Then

- **weights**: w_{jk} appears only in z_j, with $\partial z_j / \partial w_{jk} = a_k$, so

$$\frac{\partial L}{\partial w_{jk}} = dz_j\, a_k \quad \Longrightarrow \quad d\mathbf{W} = (d\mathbf{z})^\top \mathbf{a} \quad \text{(an outer product)};$$

- **activations of the previous layer**: a_k appears in **every** z_j, with $\partial z_j / \partial a_k = w_{jk}$; by the multivariable chain rule,

$$\frac{\partial L}{\partial a_k} = \sum_j dz_j\, w_{jk} \quad \Longrightarrow \quad d\mathbf{a} = d\mathbf{z}\, \mathbf{W};$$

- **through the activation**: $a_j = g(z_j)$ depends only on z_j, so $dz_j = da_j\, g'(z_j)$, i.e. $d\mathbf{z} = d\mathbf{a} \odot g'(\mathbf{z})$.

A whole mini-batch at once

With the m examples as the rows of $\mathbf{A}^{[0]} = \mathbf{X}$, and $J = \frac{1}{m}\sum_i L^{(i)}$, row i of $d\mathbf{Z}^{[l]}$ holds $d\mathbf{z}^{[l]}$ of example i:

$$d\mathbf{Z}^{[2]} = \mathbf{A}^{[2]} - \mathbf{y}, \qquad \frac{\partial J}{\partial \mathbf{W}^{[2]}} = \frac{1}{m}\left(d\mathbf{Z}^{[2]}\right)^{\top}\mathbf{A}^{[1]}, \qquad \frac{\partial J}{\partial \mathbf{b}^{[2]}} = \frac{1}{m}\sum_{i=1}^{m}\left(\text{row } i \text{ of } d\mathbf{Z}^{[2]}\right),$$

$$d\mathbf{A}^{[1]} = d\mathbf{Z}^{[2]}\,\mathbf{W}^{[2]}, \qquad d\mathbf{Z}^{[1]} = d\mathbf{A}^{[1]} \odot g'(\mathbf{Z}^{[1]}),$$

$$\frac{\partial J}{\partial \mathbf{W}^{[1]}} = \frac{1}{m}\left(d\mathbf{Z}^{[1]}\right)^{\top}\mathbf{A}^{[0]}, \qquad \frac{\partial J}{\partial \mathbf{b}^{[1]}} = \frac{1}{m}\sum_{i=1}^{m}\left(\text{row } i \text{ of } d\mathbf{Z}^{[1]}\right).$$

The factor $\frac{1}{m}$ and the sum over rows come from averaging the loss over the examples. In NumPy the sum is `dZ.sum(axis=0)`.

Backpropagation in NumPy, checked against autograd

```python
import numpy as np

def sigmoid(z):
    return 1 / (1 + np.exp(-z))

def backprop(X, y, W1, b1, W2, b2):
    m = X.shape[0]
    Z1 = X @ W1.T + b1;  A1 = sigmoid(Z1)    # forward
    Z2 = A1 @ W2.T + b2; A2 = sigmoid(Z2)
    dZ2 = A2 - y                              # (m, 1)
    dW2 = dZ2.T @ A1 / m
    db2 = dZ2.mean(axis=0)
    dA1 = dZ2 @ W2                            # (m, n1)
    dZ1 = dA1 * A1 * (1 - A1)                 # sigmoid'
    dW1 = dZ1.T @ X / m
    db1 = dZ1.mean(axis=0)
    return dW1, db1, dW2, db2
```

- Every line is one of the equations on the previous slides.
- Test: random data with $m = 8$, $d = 3$, $n^{[1]} = 4$. We compare with PyTorch's `loss.backward()` on the same network and the same binary cross-entropy.
- The largest difference among all gradient entries is 6×10^{-17}, i.e. rounding error.

Autograd does exactly this, for any network, automatically.

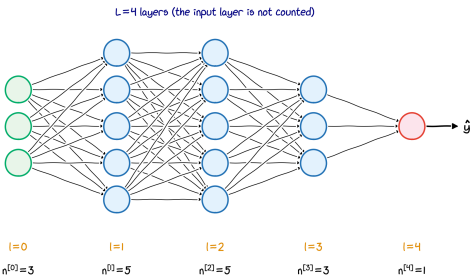
Training a network with gradient descent

1. **Initialise** all weights with small **random** values; biases may start at 0.
2. **Repeat**, for every epoch and mini-batch:
   - forward pass: compute and cache $\mathbf{Z}^{[l]}$ and $\mathbf{A}^{[l]}$ for $l = 1, \dots, L$;
   - compute the cost J;
   - backward pass: compute $\partial J / \partial \mathbf{W}^{[l]}$ and $\partial J / \partial \mathbf{b}^{[l]}$ for $l = L, \dots, 1$;
   - update: $\mathbf{W}^{[l]} \leftarrow \mathbf{W}^{[l]} - \alpha \frac{\partial J}{\partial \mathbf{W}^{[l]}}$, $\mathbf{b}^{[l]} \leftarrow \mathbf{b}^{[l]} - \alpha \frac{\partial J}{\partial \mathbf{b}^{[l]}}$.

Why random? If all weights of a layer start equal (e.g. all zero), every unit in that layer computes the same output and receives the same gradient, so they stay identical forever: $n^{[l]}$ units then behave like one. Random values **break the symmetry**. `nn.Linear` draws its weights uniformly from $\left[-1/\sqrt{n^{[l-1]}},\ 1/\sqrt{n^{[l-1]}}\right]$.

Deep neural networks

Notation for a deep network



- L: number of layers; here $L = 4$.
- $n^{[l]}$: number of units in layer l; here $n^{[0]} = 3$, $n^{[1]} = n^{[2]} = 5$, $n^{[3]} = 3$, $n^{[4]} = 1$.
- $\mathbf{W}^{[l]} \in \mathbb{R}^{n^{[l]} \times n^{[l-1]}}$, $\mathbf{b}^{[l]} \in \mathbb{R}^{n^{[l]}}$: parameters of layer l.
- $g^{[l]}$: activation function of layer l.
- $\mathbf{a}^{[l]} = g^{[l]}(\mathbf{z}^{[l]})$: activations of layer l; $\mathbf{a}^{[0]} = \mathbf{x}$ and $\mathbf{a}^{[L]} = \hat{\mathbf{y}}$.

Forward and backward pass for layer l

Forward: input $\mathbf{A}^{[l-1]}$; output $\mathbf{A}^{[l]}$; cache $\mathbf{Z}^{[l]}$

$$\mathbf{Z}^{[l]} = \mathbf{A}^{[l-1]}\big(\mathbf{W}^{[l]}\big)^\top + \mathbf{b}^{[l]}$$
$$\mathbf{A}^{[l]} = g^{[l]}\big(\mathbf{Z}^{[l]}\big)$$

for $l = 1, 2, \dots, L$, starting from $\mathbf{A}^{[0]} = \mathbf{X}$.

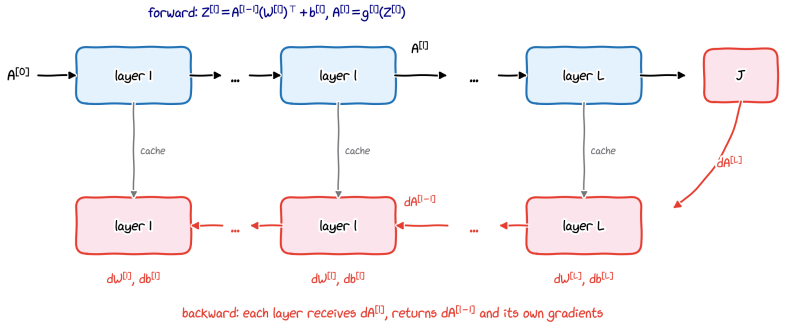
Backward: input $d\mathbf{A}^{[l]}$; output $d\mathbf{A}^{[l-1]}$ and the gradients of layer l

$$d\mathbf{Z}^{[l]} = d\mathbf{A}^{[l]} \odot g^{[l]\prime}\big(\mathbf{Z}^{[l]}\big)$$
$$\frac{\partial J}{\partial \mathbf{W}^{[l]}} = \frac{1}{m}\big(d\mathbf{Z}^{[l]}\big)^\top \mathbf{A}^{[l-1]}$$
$$\frac{\partial J}{\partial \mathbf{b}^{[l]}} = \frac{1}{m}\,\texttt{dZ.sum(axis=0)}$$
$$d\mathbf{A}^{[l-1]} = d\mathbf{Z}^{[l]}\,\mathbf{W}^{[l]}$$

for $l = L, \dots, 1$.

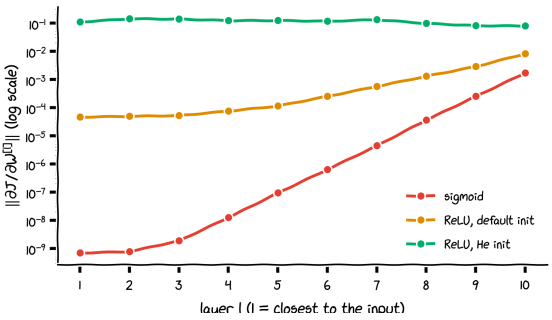
- The backward pass needs $\mathbf{A}^{[l-1]}$ and $\mathbf{Z}^{[l]}$ from the forward pass, so they are **cached**.
- Start of the backward pass: for a sigmoid output with binary cross-entropy, or softmax with cross-entropy (Lecture 3b), $d\mathbf{Z}^{[L]} = \mathbf{A}^{[L]} - \mathbf{Y}$ (with $\mathbf{Y}$ one-hot for softmax).

The whole picture



- Every layer is a building block with a forward function and a backward function.
- Deep networks just chain more of the same blocks; nothing new is needed in the mathematics. This is exactly how PyTorch's modules work.

Vanishing gradients



A network with 10 hidden layers of 64 units, at initialisation. Gradient size per layer:

- **sigmoid**: each backward step multiplies by $\sigma'(z) \leq \frac{1}{4}$. The first layer's gradient is about **a million times** smaller than the last layer's, so the first layers barely learn.
- **ReLU**, default initialisation: about 180 times smaller.
- **ReLU with He initialisation** (weights with standard deviation $\sqrt{2/n^{[l-1]}}$): the same size in every layer.

ReLU and careful initialisation are what make deep networks trainable.

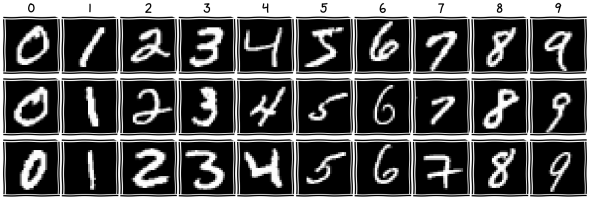
Neural networks in PyTorch

Building blocks

PyTorch	what it does
`nn.Module`	base class of every model: define the layers in `__init__`, the computation in `forward`
`nn.Linear(n_in, n_out)`	$\mathbf{Z} = \mathbf{A}\mathbf{W}^\top + \mathbf{b}$, with $\mathbf{W}$ of shape `(n_out, n_in)` and its own initialisation
`nn.ReLU()`, `nn.Sigmoid()`, `nn.Tanh()`	activation functions as layers (no parameters)
`torch.relu`, `torch.sigmoid`, `torch.tanh`	the same as functions (also in `torch.nn.functional`)
`nn.Flatten()`	turns each 28×28 image into a vector of length 784
`nn.Sequential(...)`	applies layers one after another
`nn.CrossEntropyLoss()`	softmax and cross-entropy in one step, from raw scores ("logits")

Documentation: https://pytorch.org/docs/stable/nn.html

A network for handwritten digits



$$\text{Linear}[784, 512] \to \text{ReLU} \to \text{Linear}[512, 512] \to \text{ReLU} \to \text{Linear}[512, 10]$$

Questions: Which components have learnable parameters? How many parameters does the model have?

Only the linear layers: $784 \cdot 512 + 512 = 401\,920$, $512 \cdot 512 + 512 = 262\,656$, $512 \cdot 10 + 10 = 5\,130$. In total **669,706** parameters. The 10 outputs are the scores of the digits 0 to 9.

The model in code

```
import torch
from torch import nn

class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()                    # (batch, 1, 28, 28) -> (batch, 784)
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28 * 28, 512), nn.ReLU(),
            nn.Linear(512, 512), nn.ReLU(),
            nn.Linear(512, 10),                        # 10 logits, one per digit
        )

    def forward(self, x):
        x = self.flatten(x)
        return self.linear_relu_stack(x)

device = "cuda" if torch.cuda.is_available() else "cpu"
model = NeuralNetwork().to(device)
sum(p.numel() for p in model.parameters())          # 669706
```

We only write the **forward** pass; autograd derives the backward pass.

Loading MNIST and training

```
from torch.utils.data import DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor

train_data = datasets.MNIST("data", train=True, download=True, transform=ToTensor())
test_data = datasets.MNIST("data", train=False, download=True, transform=ToTensor())
train_loader = DataLoader(train_data, batch_size=64, shuffle=True)
test_loader = DataLoader(test_data, batch_size=64)

loss_fn = nn.CrossEntropyLoss()                    # softmax + cross-entropy, from logits
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)

for epoch in range(10):
    train_loop(train_loader, model, loss_fn, optimizer)       # next slide
    test_loss, test_acc = test_loop(test_loader, model, loss_fn)
    print(f"epoch {epoch + 1}: test loss {test_loss:.3f}, accuracy {test_acc:.1%}")
```

- `ToTensor()` turns each image into a tensor of shape (1, 28, 28) with values in $[0, 1]$.
- The data are downloaded once into the folder `data`.

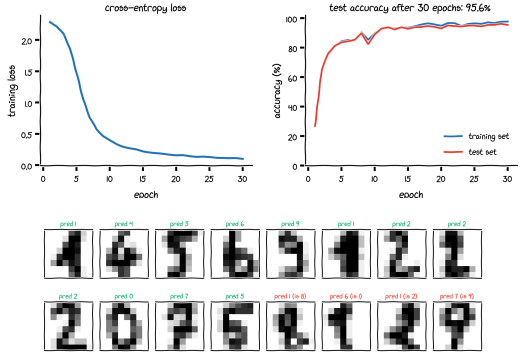
The training and test loops

```python
def train_loop(loader, model, loss_fn, optimizer):
    model.train()
    for X, y in loader:
        loss = loss_fn(model(X), y)       # forward pass and loss
        optimizer.zero_grad()
        loss.backward()                   # backpropagation
        optimizer.step()                  # gradient-descent update

def test_loop(loader, model, loss_fn):
    model.eval()
    correct, total_loss = 0, 0.0
    with torch.no_grad():                 # no gradients needed for evaluation
        for X, y in loader:
            logits = model(X)
            total_loss += loss_fn(logits, y).item() * len(y)
            correct += (logits.argmax(dim=1) == y).sum().item()
    n = len(loader.dataset)
    return total_loss / n, correct / n    # average loss, accuracy
```

`model.train()` and `model.eval()` switch layers such as dropout between training and evaluation mode; for our network they change nothing, but it is good practice.

Results



The same code on scikit-learn's 8×8 digits (1,797 images; network $64 \to 128 \to 128 \to 10$): **95.6%** test accuracy after 30 epochs; 24 of 540 test images are wrong (red). For MNIST, run the course notebook.

Summary of part 7b

- **Computational graph**: a forward pass computes values; a backward pass applies the chain rule from the output to every parameter, reusing each intermediate derivative.
- **Backpropagation**, layer l (mini-batch): $d\mathbf{Z}^{[l]} = d\mathbf{A}^{[l]} \odot g^{[l]\prime}(\mathbf{Z}^{[l]})$, $\frac{\partial J}{\partial \mathbf{W}^{[l]}} = \frac{1}{m}(d\mathbf{Z}^{[l]})^\top \mathbf{A}^{[l-1]}$, $\frac{\partial J}{\partial \mathbf{b}^{[l]}} = \frac{1}{m}\sum_{\text{rows}} d\mathbf{Z}^{[l]}$, $d\mathbf{A}^{[l-1]} = d\mathbf{Z}^{[l]}\mathbf{W}^{[l]}$.
- It starts from $d\mathbf{Z}^{[L]} = \mathbf{A}^{[L]} - \mathbf{Y}$ for sigmoid or softmax outputs with cross-entropy.
- **Initialise randomly** to break symmetry; use **ReLU** and suitable initialisation against **vanishing gradients**.
- **PyTorch**: define the layers in `nn.Module` and write `forward`; `loss.backward()` performs backpropagation; the optimiser updates the parameters.

Next, part 7c: where the backpropagation equations come from, derived step by step and checked on a small example.

Check your understanding: part 7b

1. For $J = 3(a + bc)$ at $a = 1$, $b = 2$, $c = 4$, compute $\partial J/\partial b$ and $\partial J/\partial c$.
2. In layer l, $\mathbf{W}^{[l]}$ has shape $(64, 128)$ and the mini-batch has $m = 32$ examples. What are the shapes of $\mathbf{A}^{[l-1]}$, $d\mathbf{Z}^{[l]}$ and $\partial J/\partial \mathbf{W}^{[l]}$?
3. Why do the backward formulas need $\mathbf{Z}^{[l]}$ and $\mathbf{A}^{[l-1]}$ from the forward pass?
4. What goes wrong if all weights are initialised to 0?
5. Why do gradients in early layers of a deep sigmoid network become very small?

Part 7c: the mathematics behind backpropagation

Where we are

Part 7b **stated** the backpropagation equations, checked them against autograd, and used them to train networks:

$$d\mathbf{z}^{[l]} = d\mathbf{a}^{[l]} \odot g^{[l]\prime}(\mathbf{z}^{[l]}), \quad d\mathbf{W}^{[l]} = \left(d\mathbf{z}^{[l]}\right)^{\top} \mathbf{a}^{[l-1]}, \quad d\mathbf{b}^{[l]} = d\mathbf{z}^{[l]}, \quad d\mathbf{a}^{[l-1]} = d\mathbf{z}^{[l]}\, \mathbf{W}^{[l]}.$$

Part 7c shows **why** they hold:

- the chain rule, applied entry by entry to vectors and matrices;
- every gradient as a product of **local** ("easy") derivatives and **accumulated** derivatives from the loss;
- a small network computed completely **by hand**, and checked with finite differences.

The chain rule for vectors and matrices

Vectors and matrices: the chain rule entry by entry

The chain rule of part 7b, for numbers, extends to vectors and matrices. In a network, the intermediate quantities are **vectors** ($\mathbf{z}$, $\mathbf{a}$) and the parameters are **matrices** ($\mathbf{W}$). Apply the rule to every entry, summing over every entry it passes through:

$$\frac{\partial L}{\partial W_{jk}} = \sum_i \frac{\partial L}{\partial z_i} \frac{\partial z_i}{\partial W_{jk}}.$$

We usually write this in a compact "reduced" form,

$$\frac{\partial L}{\partial \mathbf{W}} = \frac{\partial L}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{W}},$$

keeping in mind that it stands for the sum above, for every entry W_{jk}. The compact form tells us **what** to multiply; the entry-wise form tells us **how**.

One hidden layer: the derivation

The network

One example, as a row vector (as in parts 7a and 7b), hidden activation g and sigmoid output:

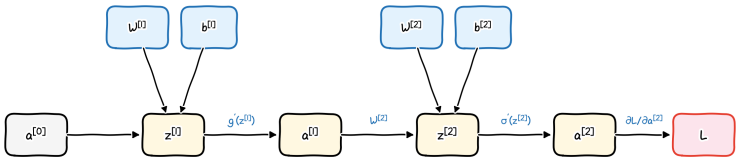
$$\mathbf{z}^{[1]} = \mathbf{a}^{[0]}\big(\mathbf{W}^{[1]}\big)^\top + \mathbf{b}^{[1]}, \quad \mathbf{a}^{[1]} = g(\mathbf{z}^{[1]}), \quad \text{with } \mathbf{a}^{[0]} = \mathbf{x},$$
$$z^{[2]} = \mathbf{a}^{[1]}\big(\mathbf{W}^{[2]}\big)^\top + b^{[2]}, \quad a^{[2]} = \sigma(z^{[2]}), \quad L = -y \log a^{[2]} - (1-y)\log\left(1 - a^{[2]}\right).$$

A chain of simple mappings:

$$\mathbf{a}^{[0]} \xrightarrow{\text{linear}} \mathbf{z}^{[1]} \xrightarrow{\text{activation}} \mathbf{a}^{[1]} \xrightarrow{\text{linear}} z^{[2]} \xrightarrow{\text{activation}} a^{[2]} \xrightarrow{\text{loss}} L(a^{[2]}, y)$$

Goal: compute $\dfrac{\partial L}{\partial \mathbf{W}^{[1]}}$, $\dfrac{\partial L}{\partial \mathbf{b}^{[1]}}$, $\dfrac{\partial L}{\partial \mathbf{W}^{[2]}}$, $\dfrac{\partial L}{\partial b^{[2]}}$.

Local and accumulated derivatives



forward: a chain of simple mappings (blue: the local, 'easy' derivative of each step)

- **Local (“easy”)** derivatives belong to a single mapping and only need the values of the forward pass: $\partial a/\partial z$, $\partial z/\partial a$, $\partial z/\partial \mathbf{W}$, $\partial z/\partial \mathbf{b}$.
- **Accumulated** derivatives $\partial L/\partial(\cdot)$ measure the effect on the **loss**. They are the hard part: each depends on everything that happens later in the chain.

The four easy derivatives

For a linear map $z_j = \sum_k w_{jk}\, a_k + b_j$ and an element-wise activation $a_j = g(z_j)$:

local derivative	value	in words
$\dfrac{\partial a_j}{\partial z_i}$	$g'(z_j)$ if $i = j$, else 0	each a_j depends only on its own z_j
$\dfrac{\partial z_j}{\partial a_k}$	w_{jk}	the weight from input k to unit j
$\dfrac{\partial z_i}{\partial w_{jk}}$	a_k if $i = j$, else 0	w_{jk} only enters z_j
$\dfrac{\partial z_i}{\partial b_j}$	1 if $i = j$, else 0	b_j only enters z_j

Activation derivatives: sigmoid $\sigma'(z) = \sigma(z)(1 - \sigma(z))$; $\tanh'(z) = 1 - \tanh^2(z)$; $\mathrm{ReLU}'(z) = 1$ for $z > 0$ and 0 for $z < 0$.

Where the backward pass starts

The loss depends on the network only through the output $a^{[2]}$. This first accumulated derivative is easy:

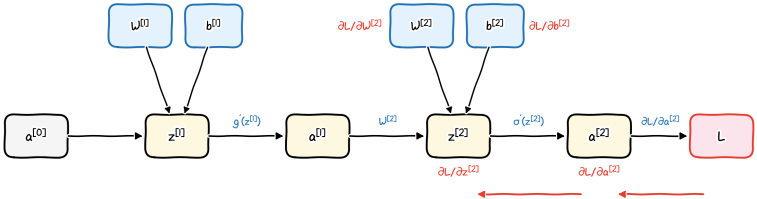
$$\frac{\partial L}{\partial a^{[2]}} = -\frac{y}{a^{[2]}} + \frac{1-y}{1-a^{[2]}}.$$

One step further back, through the sigmoid:

$$\frac{\partial L}{\partial z^{[2]}} = \frac{\partial L}{\partial a^{[2]}}\,\sigma'(z^{[2]}) = \left(-\frac{y}{a^{[2]}} + \frac{1-y}{1-a^{[2]}}\right) a^{[2]}(1-a^{[2]}) = a^{[2]} - y.$$

The prediction error, exactly as for logistic regression (Lecture 3a).

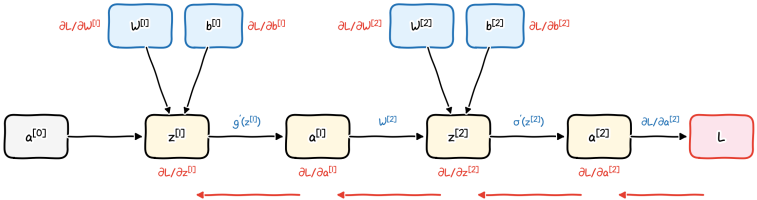
Layer 2: two steps back



layer 2: two steps back from the loss

$$\frac{\partial L}{\partial w_{1k}^{[2]}} = \frac{\partial L}{\partial z^{[2]}} \underbrace{\frac{\partial z^{[2]}}{\partial w_{1k}^{[2]}}}_{a_k^{[1]}} \implies \frac{\partial L}{\partial \mathbf{W}^{[2]}} = (a^{[2]} - y)\,\mathbf{a}^{[1]}, \qquad \frac{\partial L}{\partial b^{[2]}} = a^{[2]} - y.$$

Layer 1: continue one layer further



layer 1: the same steps, continued one layer further

$$\underbrace{\frac{\partial L}{\partial a_k^{[1]}} = \frac{\partial L}{\partial z^{[2]}}\, w_{1k}^{[2]}}_{\text{through } \partial z/\partial a} \qquad \underbrace{\frac{\partial L}{\partial z_j^{[1]}} = \frac{\partial L}{\partial a_j^{[1]}}\, g'(z_j^{[1]})}_{\text{through } \partial a/\partial z} \qquad \underbrace{\frac{\partial L}{\partial w_{jk}^{[1]}} = \frac{\partial L}{\partial z_j^{[1]}}\, a_k^{[0]}, \quad \frac{\partial L}{\partial b_j^{[1]}} = \frac{\partial L}{\partial z_j^{[1]}}}_{\text{through } \partial z/\partial \theta}$$

The backpropagation process

Order of computation, with Θ standing for $\mathbf{W}$ or $\mathbf{b}$:

$$\frac{\partial L}{\partial a^{[2]}} \xrightarrow{\partial a^{[2]}/\partial z^{[2]}} \frac{\partial L}{\partial z^{[2]}} \xrightarrow{\partial z^{[2]}/\partial \mathbf{a}^{[1]}} \frac{\partial L}{\partial \mathbf{a}^{[1]}} \xrightarrow{\partial \mathbf{a}^{[1]}/\partial \mathbf{z}^{[1]}} \frac{\partial L}{\partial \mathbf{z}^{[1]}} \xrightarrow{\partial \mathbf{z}^{[1]}/\partial \Theta^{[1]}} \frac{\partial L}{\partial \Theta^{[1]}}$$

$$\downarrow \partial z^{[2]}/\partial \Theta^{[2]}$$

$$\frac{\partial L}{\partial \Theta^{[2]}}$$

- Each red derivative is computed **once** and used twice: for the parameters of its layer, and to go one layer further back.
- Computing $\partial L/\partial \Theta^{[1]}$ is the **same process** as for $\partial L/\partial \Theta^{[2]}$, continued one layer further.
- Every arrow multiplies by one easy derivative. Nothing is ever recomputed.

The equations, collected

One hidden layer, one example ($d\theta := \partial L/\partial\theta$; $\odot$: element-wise product)

$$dz^{[2]} = a^{[2]} - y, \qquad d\mathbf{W}^{[2]} = (dz^{[2]})^\top \mathbf{a}^{[1]}, \qquad db^{[2]} = dz^{[2]},$$
$$d\mathbf{a}^{[1]} = dz^{[2]}\,\mathbf{W}^{[2]}, \qquad d\mathbf{z}^{[1]} = d\mathbf{a}^{[1]} \odot g'(\mathbf{z}^{[1]}),$$
$$d\mathbf{W}^{[1]} = (d\mathbf{z}^{[1]})^\top \mathbf{a}^{[0]}, \qquad d\mathbf{b}^{[1]} = d\mathbf{z}^{[1]}.$$

These are the entry-wise results of the last two slides, written with matrices:

- $\frac{\partial L}{\partial w_{jk}} = dz_j\, a_k$ for all j, k is the outer product $(d\mathbf{z})^\top \mathbf{a}$;
- $\frac{\partial L}{\partial a_k} = \sum_j dz_j\, w_{jk}$ is the vector-matrix product $d\mathbf{z}\,\mathbf{W}$.

A worked example

Forward pass with numbers

A $2 \to 2 \to 1$ network with sigmoid activations, one example $\mathbf{x} = (1, 2)$ with label $y = 1$:

$$\mathbf{W}^{[1]} = \begin{pmatrix} 0.1 & 0.2 \\ -0.3 & 0.4 \end{pmatrix}, \quad \mathbf{b}^{[1]} = (0,\ 0.1), \quad \mathbf{W}^{[2]} = (0.5,\ -0.6), \quad b^{[2]} = 0.2.$$

quantity	computation	value
$\mathbf{z}^{[1]}$	$(1 \cdot 0.1 + 2 \cdot 0.2,\ \ 1 \cdot (-0.3) + 2 \cdot 0.4 + 0.1)$	$(0.5,\ 0.6)$
$\mathbf{a}^{[1]}$	$\sigma(\mathbf{z}^{[1]})$	$(0.6225,\ 0.6457)$
$z^{[2]}$	$0.5 \cdot 0.6225 - 0.6 \cdot 0.6457 + 0.2$	0.1238
$a^{[2]}$	$\sigma(z^{[2]})$	0.5309
L	$-\log a^{[2]}$ (since $y = 1$)	0.6331

Backward pass with numbers

derivative	computation	value
$dz^{[2]}$	$a^{[2]} - y$	-0.4691
$d\mathbf{W}^{[2]}$	$dz^{[2]}\,\mathbf{a}^{[1]}$	$(-0.2920,\ -0.3029)$
$db^{[2]}$	$dz^{[2]}$	-0.4691
$d\mathbf{a}^{[1]}$	$dz^{[2]}\,\mathbf{W}^{[2]} =$ $-0.4691 \cdot (0.5,\ -0.6)$	$(-0.2345,\ 0.2814)$
$g'(\mathbf{z}^{[1]})$	$\mathbf{a}^{[1]} \odot (1 - \mathbf{a}^{[1]})$	$(0.2350,\ 0.2288)$
$d\mathbf{z}^{[1]}$	$d\mathbf{a}^{[1]} \odot g'(\mathbf{z}^{[1]})$	$(-0.0551,\ 0.0644)$
$d\mathbf{W}^{[1]}$	$(d\mathbf{z}^{[1]})^\top \mathbf{a}^{[0]}$	$\begin{pmatrix} -0.0551 & -0.1102 \\ 0.0644 & 0.1288 \end{pmatrix}$
$d\mathbf{b}^{[1]}$	$d\mathbf{z}^{[1]}$	$(-0.0551,\ 0.0644)$

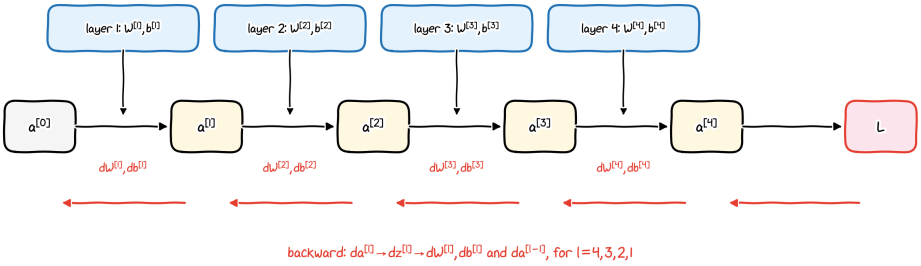
Check: changing each parameter by $\pm 10^{-6}$ and recomputing L (finite differences) gives the same nine gradient entries, to within 10^{-10}.

Reading the numbers

- $dz^{[2]} = -0.47 < 0$: the network predicts 0.53 for a positive example, so increasing $z^{[2]}$ would reduce the loss. Gradient descent ($\theta \leftarrow \theta - \alpha\, d\theta$) does exactly that.
- $d\mathbf{W}^{[2]} < 0$: both hidden units are active ($a^{[1]} \approx 0.6$), so increasing either weight increases $z^{[2]}$.
- Hidden unit 1 has a **positive** outgoing weight (0.5) and unit 2 a **negative** one (-0.6): so $d\mathbf{z}^{[1]}$ has opposite signs. Training pushes unit 1 to be more active and unit 2 to be less active.
- The gradients of layer 1 are 2 to 5 times smaller than those of layer 2: each step back through a sigmoid multiplies by $\sigma' \leq 0.25$ (vanishing gradients, part 7b).

Any number of layers

Four layers: the same step, four times



$$\frac{\partial L}{\partial \mathbf{a}^{[4]}} \to \frac{\partial L}{\partial \mathbf{z}^{[4]}} \to \frac{\partial L}{\partial \mathbf{a}^{[3]}} \to \frac{\partial L}{\partial \mathbf{z}^{[3]}} \to \frac{\partial L}{\partial \mathbf{a}^{[2]}} \to \frac{\partial L}{\partial \mathbf{z}^{[2]}} \to \frac{\partial L}{\partial \mathbf{a}^{[1]}} \to \frac{\partial L}{\partial \mathbf{z}^{[1]}}$$

and from each $\partial L/\partial \mathbf{z}^{[l]}$ we branch off to $\partial L/\partial \Theta^{[l]}$ using $\partial \mathbf{z}^{[l]}/\partial \Theta^{[l]}$.

To get the gradient of layer $l-1$, go through **exactly the same process** as for layer l, then continue one layer further. The four equations of this step are those of part 7b ("Forward and backward pass for layer l").

The algorithm

```
forward pass:   A[0] = X
                for l = 1, ..., L:
                    Z[l] = A[l-1] @ W[l].T + b[l]          # cache Z[l] and A[l-1]
                    A[l] = g[l](Z[l])

backward pass:  dZ = A[L] - Y                              # output layer, cross-entropy loss
                for l = L, ..., 1:
                    dW[l] = dZ.T @ A[l-1] / m
                    db[l] = dZ.sum(axis=0) / m
                    if l > 1:
                        dA = dZ @ W[l]                     # pass the error back
                        dZ = dA * g'[l-1](Z[l-1])          # through the activation

update:         W[l] -= alpha * dW[l];   b[l] -= alpha * db[l]     for every l
```

Each layer needs only its own cached values and the error passed back from the layer above.

Summary of part 7c

- **Chain rule**: multiply local derivatives along a path; sum over all paths. For vectors and matrices, apply it entry by entry.
- **Local ("easy") derivatives**: $\partial \mathbf{a}/\partial \mathbf{z} = g'(\mathbf{z})$ (element-wise), $\partial \mathbf{z}/\partial \mathbf{a} = \mathbf{W}$, $\partial z_j/\partial w_{jk} = a_k$, $\partial z_j/\partial b_j = 1$.
- **Accumulated derivatives** $\partial L/\partial(\cdot)$ are passed backwards, computed once and reused.
- The output layer starts with $dz^{[L]} = a^{[L]} - y$; every layer then applies the same four equations.
- Worked example: hand computation and finite differences agree.

The optional appendix to this lecture gives a compact, rigorous derivation of the same results with matrix calculus.

Check your understanding: part 7c

1. With $f = a_1 + a_1 a_2$, $a_1 = 2\theta$, $a_2 = \theta^2$: compute $\partial f / \partial \theta$ with the chain rule.
2. Why is $\partial a_j / \partial z_i = 0$ for $i \neq j$, but $\partial z_j / \partial a_k \neq 0$ for all k?
3. In the worked example, compute $dW_{11}^{[2]}$ yourself from $dz^{[2]}$ and $\mathbf{a}^{[1]}$.
4. Which values from the forward pass does the backward step for layer l need?
5. What changes in the backward pass if the hidden activation is ReLU instead of sigmoid?

Wrapping up

Lecture 7 at a glance

	one layer l, mini-batch (rows = examples)
forward	$\mathbf{Z}^{[l]} = \mathbf{A}^{[l-1]}(\mathbf{W}^{[l]})^\top + \mathbf{b}^{[l]}$, $\mathbf{A}^{[l]} = g^{[l]}(\mathbf{Z}^{[l]})$
start of backward	$d\mathbf{Z}^{[L]} = \mathbf{A}^{[L]} - \mathbf{Y}$ (sigmoid or softmax with cross-entropy)
backward	$d\mathbf{Z}^{[l]} = d\mathbf{A}^{[l]} \odot g^{[l]\prime}(\mathbf{Z}^{[l]})$, $d\mathbf{A}^{[l-1]} = d\mathbf{Z}^{[l]}\mathbf{W}^{[l]}$
gradients	$\frac{\partial J}{\partial \mathbf{W}^{[l]}} = \frac{1}{m}(d\mathbf{Z}^{[l]})^\top \mathbf{A}^{[l-1]}$, $\frac{\partial J}{\partial \mathbf{b}^{[l]}} = \frac{1}{m}\sum_{\text{rows}} d\mathbf{Z}^{[l]}$
in PyTorch	`nn.Sequential(nn.Linear(...), nn.ReLU(), ...)`; `loss.backward()`; `optimizer.step()`
in practice	random initialisation, ReLU, suitable scaling against vanishing gradients

Next lecture: convolutional neural networks, which exploit the structure of images.